

Project Part 1

Pattern Recognition

ECE 759

Kudiyar Orazymbetov (korazym@ncsu.edu)
Nico Casale (ncasale@ncsu.edu)

March 14, 2018

Contents *(Note that the entries are links.)*

List of Figures	1	3	Algorithm Implementations	3
1 Introduction	2		3.1 Decision Tree Algorithm	3
			3.1.1 Entropy and Information Gain . .	6
2 Feature Selection	2		3.2 Linear Discriminant Analysis	7
2.1 Decision Tree Feature Generation	2		3.2.1 Helpful Functions	8
2.1.1 MNIST	2			
2.2 LDA vs PCA	2	4	References	8
2.3 LDA for MNIST and Extended YaleB . .	3	5	Code Listings	8

List of Figures

Listings

1 Introduction

2 Feature Selection

2.1 Decision Tree Feature Generation

2.1.1 MNIST

In working with the decision trees, we utilized the SVD of each image in the training set.

2.2 LDA vs PCA

LDA and PCA linearly transforms the data to reduce the dimensions. The difference is LDA is supervised whereas PCA is unsupervised. PCA finds directions which maximize the variance along that direction and they are orthogonal to each other. We can view a PCA technique as in Figure(2.1).




Figure 2.1: PCA implemented on 2-dimensional data

LDA maximizes the class separability and can be viewed as in Figure(2.2).




Figure 2.2: LDA visualization

2.3 LDA for MNIST and Extended YaleB

LDA reduces the data with K classes with large dimensions to $K-1$ dimensions which captures the most energy of the data. Implementing LDA as a dimension reduction, MNIST and Extended YaleB datasets are reduced to 9 and 37 variables for each data point. Then we worked with this transformed data to do further classification.

3 Algorithm Implementations

3.1 Decision Tree Algorithm

A binary decision tree is a hierarchical structure that takes input data at its root and propagates it to one of many leaves. Each *leaf* of the tree represents a class designation. To reach a leaf, the features of the data are utilized at *nodes* to make a binary decision: to proceed down the left or right *branch* of the tree? To answer this question, the node also carries a *threshold* that the feature value of the test data is compared against. If the test feature is less than the threshold, we proceed down the left branch. Otherwise, the right. An illustration of a simple decision tree is pictured below.

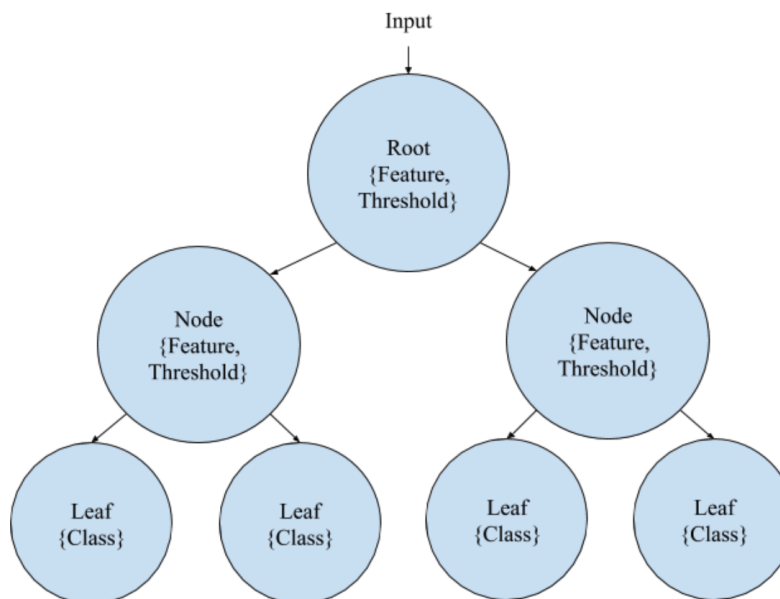


Figure 3.1: An example decision tree.

This decision tree structure needs to be generated before it can be used with test data. To train a decision tree that appropriately classifies our test data according to the features we generated, we employ a recursive function. The function signature is

```
tree = trainDecisionTree(set)
```

Where **set** is the training set, which is a MATLAB structure that contains the raw data (unused), class labels, and generated features. **tree** is the returned structure that can be used during testing. It is essentially a nested structure that contains two types of elements: nodes and leaves. At each node of the tree, a feature and threshold are specified. If a test sample's value at that particular feature is less than the threshold, the sample is passed down the left branch of the node. Similarly, if the sample's feature is greater than the threshold, it goes through the right branch. This is repeated until we reach a leaf node, which specifies a class membership.

The decision tree training algorithm has a few major steps, and proceeds by evaluating a metric called *information gain* at various configurations. For now, suffice it to say that information gain is a scalar that represents the improvement in prediction as we narrow down the set (by growing the tree) to find appropriate leaves. Our implementation most closely follows that of the ID3 and C4.5 algorithms developed by Ross Quinlan in the late 80s and early 90s [1] [2].

1. Check stopping conditions, which generate leaves.
 - If there are no more features to split on, return a leaf with the class mode of the set.
 - The set is smaller than **minLeaf**, which is a tuning parameter that is meant to reduce overfitting of the training data. If this condition is met, return a leaf with the class mode of the set.
 - If all samples in the set belong to the same class, return a leaf with the class.
 - If no feature yields an improvement to the information gain (discussed below), then return a leaf with the class mode of the set. Note that this condition is only evaluated after step 2.
2. Iterate over each feature. Sort the set along the current feature. We utilize a threshold that splits the set between adjacent feature values. Because the information gain across thresholds is convex on the whole (see Fig. 3.2), we use a line search that approximates the highest information gain for each threshold.

Let **attributeBest** and **indBest** be the feature and index that yield the highest information gain. Since the set is sorted, we can simply split the set at the index given by **indBest** for the recursion.

3. Recur over the subsets given by `indBest` to find the next attribute that yields the highest information gain. Note that we exclude the attribute/feature we chose in this execution of `trainDecisionTree(.)`.

The algorithm is reproduced in pseudocode below.

Algorithm 3.1: `trainDecisionTree`

Data: *set* of training samples with class labels (*set*(1)) and attributes (features) (*set*(2)).

minLeaf, an integer specifying the minimum number of elements in *set* required to make a splitting node. *i.e.* a leaf is generated if the set has fewer elements than *minLeaf*.

Result: *tree*, a structure containing nodes and leaves.

```

1 begin
2   check for base cases:
3   if set(2) =  $\emptyset$  then
4     no more attributes (features) to split on.
5     return leaf with mode of set(1) (class labels)
6   if length(set(1)) < minLeaf then
7     return leaf with mode of set(1) (class labels)
8   thisSetEntropy  $\leftarrow$  getEntropy(set)
9   if not thisSetEntropy then
10    all samples are in the same class.
11    return leaf with first element of set(1)
12
13   instantiate tracking variables:
14   attributeBest  $\leftarrow$  0
15   thresholdBest  $\leftarrow$  0
16   infoGainBest  $\leftarrow$  0
17   indexBest  $\leftarrow$  0
18   for attribute  $\in$  range(#offeaturesleftinset) do
19     sort set along attribute.
20     perform line search heuristic to approximate max information gain by splitting set at various indices.
21
22     if thisInfoGain > infoGainBest then
23       attributeBest  $\leftarrow$  attribute
24       infoGainBest  $\leftarrow$  thisInfoGain
25       indexBest  $\leftarrow$  index given by line search
26       thresholdBest  $\leftarrow$  attribute value midway between indexBest and indexBest + 1
27
28   if not infoGainBest or not attributeBest then
29     no attribute provides an information gain.
30     return leaf with mode of set(1) (class labels)
31
32   re-sort set along attributeBest.
33   subsets  $\leftarrow$  getSubsets(set, attributeBest, indexBest)
34   subtree1  $\leftarrow$  trainDecisionTree(subsets(1), minLeaf)
35   subtree2  $\leftarrow$  trainDecisionTree(subsets(2), minLeaf)
36
37   return node with attributeBest, thresholdBest, subtree1, and subtree2

```

Once we have trained the decision tree, we can begin to test it by passing the features generated from the test set

through the tree. The testing algorithm is reproduced in pseudocode below.

Algorithm 3.2: testDecisionTree

Data: *set* of test samples with class labels (*set*(1)) and attributes (features) (*set*(2)).

tree, the struct of structs that represents the trained decision tree.

Result: *set*, the input structure modified to include predicted class membership.

```

1 begin
2   for each sample ∈ set do
3     treeWalked ← tree
4     classified ← false
5     attributes ← features corresponding to this sample
6     while not classified do
7       if treeWalked is a node then
8         thisAttribute ← attribute given by treeWalked
9         thisAttributeValue ← value of the sample's feature at thisAttribute
10        thisThreshold ← threshold given by treeWalked
11        remove thisAttribute from the sample's feature vector.
12        compare thisAttributeValue to thisThreshold:
13        if thisAttributeValue < thisThreshold then
14          choose left branch.
15          treeWalked ← left branch given by current treeWalked.
16        else
17          choose right branch.
18          treeWalked ← right branch given by current treeWalked.
19      else
20        treeWalked is a leaf.
21        append predicted label to set(2) at this sample.
22        classified ← true

```

3.1.1 Entropy and Information Gain

When growing the decision tree, the training algorithm utilizes a metric called *information gain* to optimize the predictive power of the tree. Before we can define information gain, we must first understand *entropy*. Entropy is an information theoretic topic that represents the amount of uncertainty in a given set of data. It is defined as

$$H(X) = - \sum_{i=1}^n P(x_i) \log_2 P(x_i) \quad (3.1)$$

Where $P(x_i)$ is the proportion of the number of elements with class x_i to the number of elements in the set X . there are n classes in the set. Note that if all samples belong to class i , $P(x_i) = 1$ and $\log_2 P(x_i) = 0$, so the entropy is zero.

In calculating the information gain of splitting the set into two subsets, we utilize the entropy of the parent set and subtract from it a weighted entropy of each subset. This is better expressed by the following equation:

$$IG(X) = H(X) - \sum_{i=1}^2 \frac{|S_i|}{|X|} H(S_i) \quad (3.2)$$




Figure 3.2: Information Gain across all possible thresholds.

3.2 Linear Discriminant Analysis

Our classification criterion is to misclassify as small as possible. The rule we employ in classifying the data points is through the use of Bayes Rule. We put the data point to the group with the highest conditional probability (i.e. $P(w_i|x)$). In practice, it is not feasible to get conditional probability for a given point unless we have a huge data. So we should assume the distribution and calculate the probabilities.

LDA relies on the assumption of normal distribution of data for each class. Linear discriminant analysis frequently achieves good performances in the tasks of face and object recognition, even though the assumptions of common covariance matrix among groups and normality are often violated (Duda, et al., 2001) (Tao Li, et al., 2006).

Since MNIST and Yale datasets are high-dimensional, we can not check the normality of variables. Instead, we can reduce the dimensions via a projection. We need to employ multiclass LDA as we have several classes. The class separation in this case will be given by the ratio of $\frac{\mathbf{w}^T \Sigma_b \mathbf{w}}{\mathbf{w}^T \Sigma_w \mathbf{w}}$. In the case of two classes, this will reduce just to the ratio of between-class variance to within-class variance.

The steps that we follow in our LDA algorithm are:

1. We calculate within-class scatter matrix $\Sigma_i = \frac{1}{N_i - 1} \sum_{\mathbf{x} \in D_i} (\mathbf{x} - \boldsymbol{\mu}_i) (\mathbf{x} - \boldsymbol{\mu}_i)^T$ for each class, then sum them up to get $\Sigma_W = \sum_{i=1}^c (N_i - 1) \Sigma_i$.
2. Then we find average within-class scatter matrix by calculating $w_a = \frac{\Sigma_W}{\#of\ classes}$
3. We also calculate the between-class scatter matrix by $S_B = \sum_{i=1}^c N_i (\boldsymbol{\mu}_i - \boldsymbol{\mu})(\boldsymbol{\mu}_i - \boldsymbol{\mu})^T$.
4. We need to find eigenvectors and eigenvalues of $\Sigma^{-1} \Sigma_b$.

5. We then sort the eigenvectors depending on the magnitude of eigenvalues.
6. The number of linear discriminants will be $c - 1$ which will be 9. We need to verify that.
7. Project our data onto the subspace(constructed by the eigenvectors of the highest eigenvalues).
8. Since we have reduced dimension of our data, we can easily apply a discriminant function for a test vector to see which class it belongs
9. A discriminant function for each class is $f(i) = \mu_i w_a^{-1} x_k^T - \frac{1}{2} \mu_i w_a^{-1} \mu_i^T + \ln(P_i)$ where P_i is respective probabilities of each class.
10. Then we select a class with highest discriminant function.

3.2.1 Helpful Functions

To accomplish the above pseudo code, we used own created functions.

4 References

- [1] Wikipedia contributors. Id3 algorithm — wikipedia, the free encyclopedia, 2017. [Online; accessed 14-March-2018].
- [2] Wikipedia contributors. C4.5 algorithm — wikipedia, the free encyclopedia, 2018. [Online; accessed 14-March-2018].

5 Code Listings

Below are the primary scripts that solve the project. Please see the `code` folder for supporting functions.